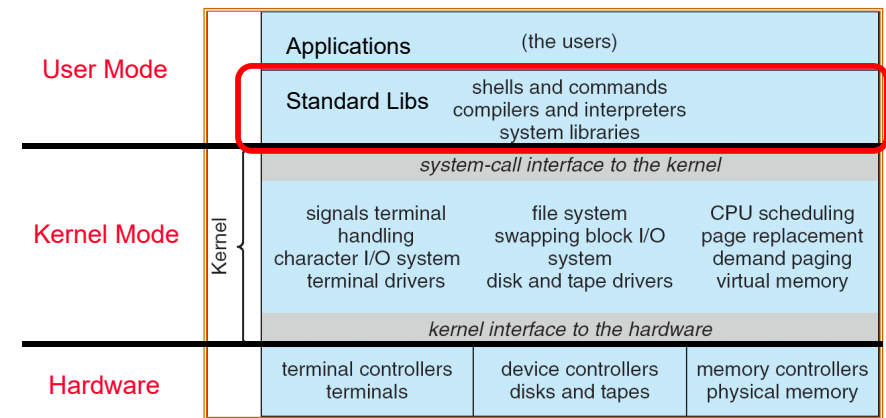


CS162 Operating Systems and Systems Programming Lecture 8

Introduction to I/O, Sockets, Networking

February 18th, 2020
Prof. John Kubiatowicz
<http://cs162.eecs.Berkeley.edu>

Recall: UNIX System Structure

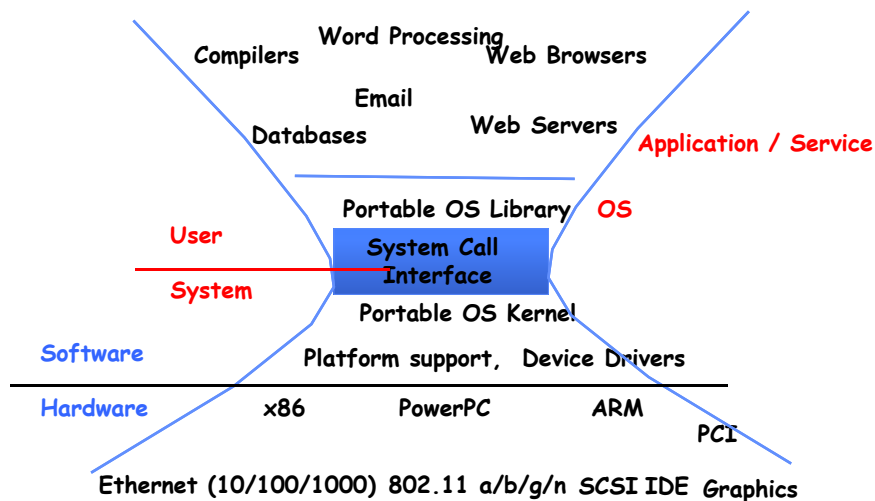


2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.2

Recall: A Kind of Narrow Waist




2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.3

Recall: web server

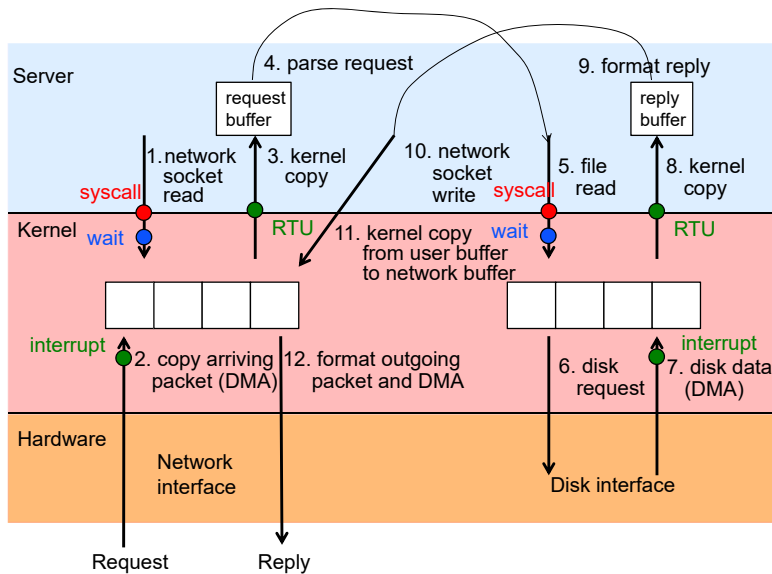


2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.4

Recall: web server



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.5

POSIX I/O: Everything is a "File"

- Identical interface for:
 - Devices (terminals, printers, etc.)
 - Regular files on disk
 - Networking (sockets)
 - Local interprocess communication (pipes, sockets)
- Based on **open()**, **read()**, **write()**, and **close()**
- Allows simple composition of programs
 - » `find | grep | wc ...`

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.6

POSIX I/O Design Patterns

- **Open before use**
 - Access control check, setup happens here
- **Byte-oriented**
 - Least common denominator
 - OS responsible for hiding the fact that real devices may not work this way (e.g. hard drive stores data in blocks)
- **Explicit close**

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.7

POSIX I/O: Kernel Buffering


- **Reads are buffered**
 - Part of making everything byte-oriented
 - Process is **blocked** while waiting for device
 - Let other processes run while gathering result
- **Writes are buffered**
 - Complete in background (more later on)
 - Return to user when data is "handed off" to kernel

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.8

Putting it together: web server

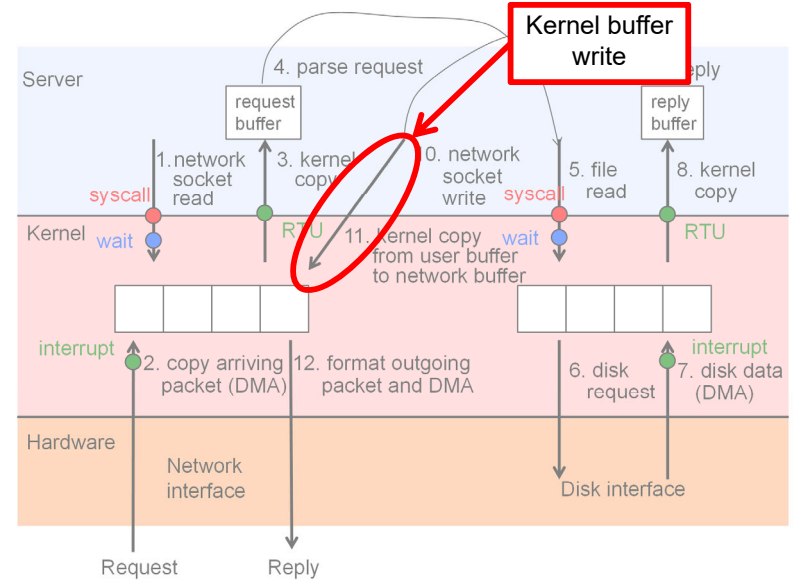


2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.9

Putting it together: web server

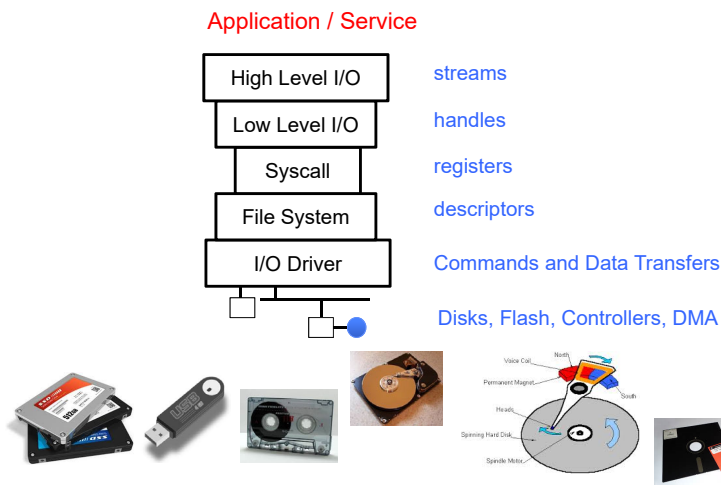


2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.10

I/O & Storage Layers



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.11

The File System Abstraction

- High-level idea
 - Files live in hierarchical namespace of filenames
- File
 - Named collection of data in a file system
 - POSIX File data: sequence of bytes
 - » Text, binary, linearized objects, ...
 - File Metadata: information about the file
 - » Size, Modification Time, Owner, Security info
 - » Basis for access control
- Directory
 - “Folder” containing files & Directories
 - Hierarchical (graphical) naming
 - » Path through the directory graph
 - » Uniquely identifies a file or directory
 - /home/ff/cs162/public_html/fa18/index.html
 - Links and Volumes (later)

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.12

C High-Level File API – Streams

- Operate on “streams” - sequence of bytes, whether text or data, with a position



```
#include <stdio.h>
FILE *fopen( const char *filename, const char *mode );
int fclose( FILE *fp );
```

Mode	Text	Binary	Descriptions
r		rb	Open existing file for reading
w		wb	Open for writing; created if does not exist
a		ab	Open for appending; created if does not exist
r+		rb+	Open existing file for reading & writing.
w+		wb+	Open for reading & writing; truncated to zero if exists, create otherwise
a+		ab+	Open for reading & writing. Created if does not exist. Read from beginning, write as append

Don't forget to flush

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.13

Connecting Processes, Filesystem, and Users

- Process has a ‘current working directory’
- Absolute Paths
 - /home/ff/cs162
- Relative paths
 - index.html, ./index.html - current WD
 - ../index.html - parent of current WD
 - ~, ~cs162 - home directory

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.14

C API Standard Streams – stdio.h

- Three predefined streams are opened implicitly when a program is executed
 - FILE *stdin – normal source of input, can be redirected
 - FILE *stdout – normal source of output, can be redirected
 - FILE *stderr – diagnostics and errors, can be redirected
- STDIN / STDOUT enable composition in Unix
- All can be redirected (for instance, using “pipe” symbol: '|’):
 - cat hello.txt | grep “World!”
 - » Cat’s stdout goes to grep’s stdin!

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.15

C high level File API – stream ops

```
#include <stdio.h>
// character oriented
int fputc(int c, FILE *fp); // rtn c or EOF on err
int fputs(const char *s, FILE *fp); // rtn >0 or EOF

int fgetc( FILE * fp );
char *fgets( char *buf, int n, FILE *fp );

// block oriented
size_t fread(void *ptr, size_t size_of_elements,
             size_t number_of_elements, FILE *a_file);

size_t fwrite(const void *ptr, size_t size_of_elements,
             size_t number_of_elements, FILE *a_file);

// formatted
int fprintf(FILE *restrict stream, const char *restrict format, ...);
int fscanf(FILE *restrict stream, const char *restrict format, ...);
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.16

C Streams: char by char I/O

```
#include <stdio.h>

int main(void) {
    FILE* input = fopen("input.txt", "r");
    FILE* output = fopen("output.txt", "w");
    int c;

    c = fgetc(input);
    while (c != EOF) {
        fputc(output, c);
        c = fgetc(input);
    }
    fclose(input);
    fclose(output);
}
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.17

Administrivia

- Lecture 7 video posted on the schedule (sorry about projector fiasco):
 - Please watch it! I'm assuming that you will master that material for the Midterm
 - I have been posting answers to questions on Piazza
 - On Thursday, I'll give some "recall" slides at the beginning and give a chance for a couple of in-class questions
- Midterm 1 is next Thursday (2/27)!
 - Please answer questions about conflicts by this Thursday
 - » New conflict time and room up now. We are trying to handle all conflicts on Thursday 2/27 instead of Friday...
 - Topics: All material up to next Tuesday
- Review Session: Tuesday (2/25)
 - Right here from 6:30-8:00pm (after class)
- Class Attendance Points:
 - Please go to alexwu.io/a and fill out the form

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.18

What if we wanted block by block I/O?

```
#include <stdio.h>
// character oriented
int fputc(int c, FILE *fp); // rtn c or EOF on err
int fputs(const char *s, FILE *fp); // rtn >0 or EOF

int fgetc( FILE *fp );
char *fgets( char *buf, int n, FILE *fp );

// block oriented
size_t fread(void *ptr, size_t size_of_elements,
             size_t number_of_elements, FILE *a_file);

size_t fwrite(const void *ptr, size_t size_of_elements,
             size_t number_of_elements, FILE *a_file);

// formatted
int fprintf(FILE *restrict stream, const char *restrict format, ...);
int fscanf(FILE *restrict stream, const char *restrict format, ...);
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.19

stdio Block-by-Block I/O

```
#include <stdio.h>
#define BUFFER_SIZE 1024
int main(void) {
    FILE* input = fopen("input.txt", "r");
    FILE* output = fopen("output.txt", "w");
    char buffer[BUFFER_SIZE];
    size_t length;
    length = fread(buffer, BUFFER_SIZE, sizeof(char), input);

    }
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.20

stdio Block-by-Block I/O

```
#include <stdio.h>
#define BUFFER_SIZE 1024
int main(void) {
    FILE* input = fopen("input.txt", "r");
    FILE* output = fopen("output.txt", "w");
    char buffer[BUFFER_SIZE];
    size_t length;
    length = fread(buffer, BUFFER_SIZE, sizeof(char), input);
    while (length > 0) {
        fwrite(buffer, length, sizeof(char), output);
        length = fread(buffer, BUFFER_SIZE, sizeof(char), input);
    }
    fclose(input);
    fclose(output);
}
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.21

Aside: Systems Programming

- Systems programmers are paranoid
- We should really be writing things like:

```
FILE* input = fopen("input.txt", "r");
if (input == NULL) {
    // Prints our string and error msg.
    perror("Failed to open input file")
}
```
- Be **thorough about checking return values**
 - Want failures to be systematically caught and dealt with

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020



Lec 8.22

C Stream API: Positioning

```
int fseek(FILE *stream, long int offset, int whence);

long int ftell (FILE *stream)

void rewind (FILE *stream)
```



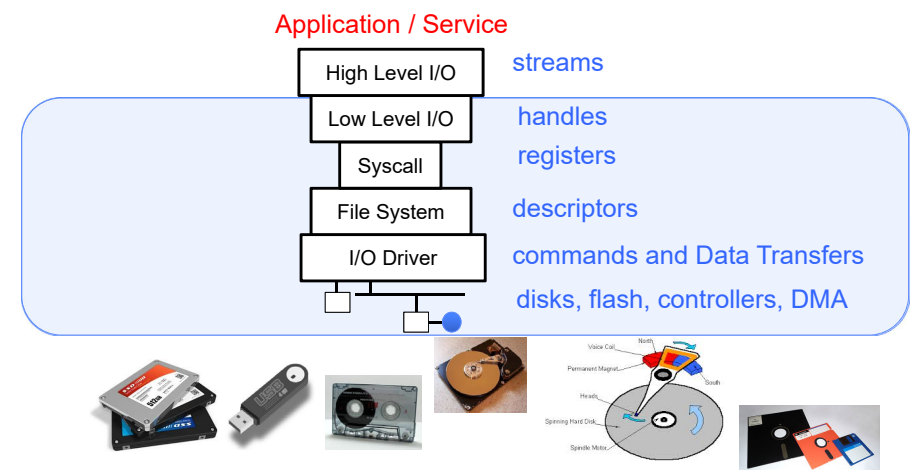
- Preserves high level abstraction of a uniform stream of objects

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.23

What's below the surface ??



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.24

C Low level I/O

- Operations on File Descriptors – as OS object representing the state of a file
 - User has a “handle” on the descriptor

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>
```

```
int open (const char *filename, int flags [, mode_t mode])
int creat (const char *filename, mode_t mode)
int close (int filedes)
```

Bit vector of:

- Access modes (Rd, Wr, ...)
- Open Flags (Create, ...)
- Operating modes (Appends, ...)

Bit vector of Permission Bits:

- User|Group|Other X R|W|X

http://www.gnu.org/software/libc/manual/html_node/Opening-and-Closing-Files.html

C Low Level: standard descriptors

```
#include <unistd.h>
```

```
STDIN_FILENO - macro has value 0
STDOUT_FILENO - macro has value 1
STDERR_FILENO - macro has value 2
```

```
int fileno (FILE *stream)
```

```
FILE * fdopen (int filedes, const char *opentype)
```

- Crossing levels: File descriptors vs. streams
- Don't mix them!

C Low Level Operations

```
ssize_t read (int filedes, void *buffer, size_t maxsize)
- returns bytes read, 0 => EOF, -1 => error
ssize_t write (int filedes, const void *buffer, size_t size)
- returns bytes written
```

```
off_t lseek (int filedes, off_t offset, int whence)
```

```
int fsync (int filedes) - wait for i/o to finish
void sync (void) - wait for ALL to finish
```

- When write returns, data is on its way to disk and can be read, but it may not actually be permanent!

A little example: lowio.c

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    char buf[1000];
    int fd = open("lowio.c", O_RDONLY, S_IRUSR | S_IWUSR);
    ssize_t rd = read(fd, buf, sizeof(buf));
    int err = close(fd);
    ssize_t wr = write(STDOUT_FILENO, buf, rd);
}
```

And lots more !

- TTYs versus files
- Memory mapped files
- File Locking
- Asynchronous I/O
- Generic I/O Control Operations
- Duplicating descriptors

```
int dup2 (int old, int new)
int dup (int old)
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.29

Another: lowio-std.c

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>

#define BUFSIZE 1024

int main(int argc, char *argv[])
{
    char buf[BUFSIZE];
    ssize_t writelen = write(STDOUT_FILENO, "I am a process.\n", 16);

    ssize_t readlen = read(STDIN_FILENO, buf, BUFSIZE);

    ssize_t strlen = snprintf(buf, BUFSIZE, "Got %zd chars\n", readlen);

    writelen = strlen < BUFSIZE ? strlen : BUFSIZE;
    write(STDOUT_FILENO, buf, writelen);

    exit(0);
}
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.30

Low-Level I/O: Example

```
#include <fcntl.h>
#include <unistd.h>

#define BUFFER_SIZE 1024

int main(void) {
    int input_fd = open("input.txt", O_RDONLY);
    int output_fd = open("output.txt", O_WRONLY);
    char buffer[BUFFER_SIZE];
    ssize_t length;
    length = read(input_fd, buffer, BUFFER_SIZE);
    while (length > 0) {
        write(output_fd, buffer, length);
        length = read(input_fd, buffer, BUFFER_SIZE);
    }
    close(input_fd);
    close(output_fd);
}
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.31

Low-Level I/O: Other Operations

- Operations specific to terminals, devices, networking, ...
- Duplicating descriptors
 - int dup2(int old, int new);
 - int dup(int old);
- Pipes - bi-directional channel
 - int pipe(int fileds[2]);
 - Writes to fileds[1] read from fileds[0]
- File Locking
- Memory-Mapping Files
- Asynchronous I/O

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.32

Little pipe example

```
#include <unistd.h>

#define BUFSIZE 1024
enum PipeSel {rd_pipe = 0, wt_pipe = 1};

int main(int argc, char *argv[])
{
    char *msg = "Message in a pipe.\n";
    char buf[BUFSIZE];
    int pipe_fd[2];
    if (pipe(pipe_fd)) {
        fprintf(stderr, "Pipe failed.\n"); return EXIT_FAILURE;
    }
    ssize_t writelen = write(pipe_fd[wt_pipe], msg, strlen(msg)+1);
    printf("Sent: %s [%ld, %ld]\n", msg, strlen(msg)+1, writelen);

    ssize_t readlen = read(pipe_fd[rd_pipe], buf, BUFSIZE);
    printf("Rcvd: %s [%ld]\n", msg, readlen);
    close(pipe_fd[wt_pipe]);
    close(pipe_fd[rd_pipe]);
}
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.33

Inter-Process Communication (IPC)

- One process reads a file the other writes, or ...

```
pid_t pid = fork();
if (pid < 0) {
    fprintf(stderr, "Fork failed.\n");
    return EXIT_FAILURE;
}
if (pid != 0) {
    ssize_t writelen = write(pipe_fd[wt_pipe], msg, msglen);
    printf("Parent: %s [%ld, %ld]\n", msg, msglen, writelen);
    close(pipe_fd[wt_pipe]);
} else {
    ssize_t readlen = read(pipe_fd[rd_pipe], buf, BUFSIZE);
    printf("Child Rcvd: %s [%ld]\n", msg, readlen);
    close(pipe_fd[rd_pipe]);
}
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.34

Streams vs. File Descriptors

- Streams are **buffered in user memory**:
printf("Beginning of line ");
sleep(10); // sleep for 10 seconds
printf("and end of line\n");
⇒ Prints out **everything at once**
- Operations on file descriptors are **visible immediately**
write(STDOUT_FILENO, "Beginning of line ", 18);
sleep(10);
write("and end of line \n", 16);
⇒ Outputs "Beginning of line" 10 seconds earlier

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.35

Why Buffer in Userspace? Overhead!

- Avoid system call overhead
 - Time to copy registers, transition to kernel mode, jump to system call handler, etc.
- Minimum syscall time: ~100s of nanoseconds
 - Read/write a file byte by byte?
 - Max throughput of ~10MB/second
 - With fgetc? Keeps up with your SSD

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.36

Why Buffer in Userspace? Functionality.

- System call operations less capable
 - Simplifies operating system
- Example: No "read until new line" operation
 - Solution: Make a big read syscall, find first new line in userspace
 - Could simulate by one syscall per character, but we already know this is a bad idea

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.37

Summary: Key Unix I/O Design Concepts

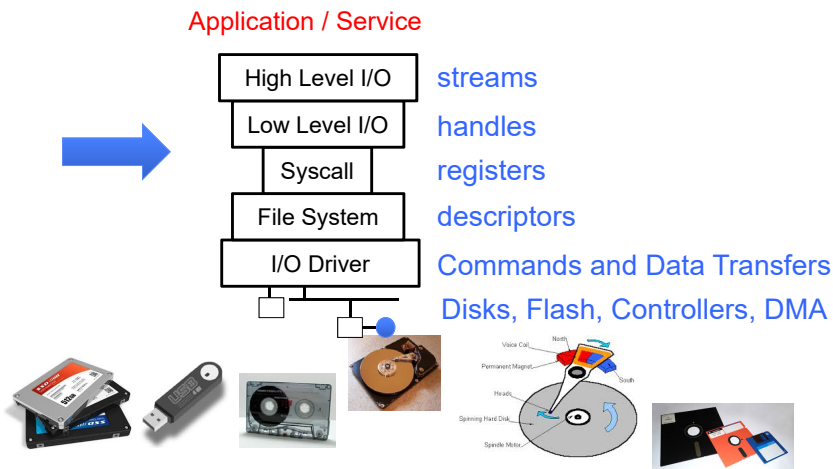
- Uniformity – everything is a file
 - file operations, device I/O, and interprocess communication through open, read/write, close
 - Allows simple composition of programs
 - » find | grep | wc ...
- Open before use
 - Provides opportunity for access control and arbitration
 - Sets up the underlying machinery, i.e., data structures
- Byte-oriented
 - Even if blocks are transferred, addressing is in bytes
- Kernel buffered reads
 - Streaming and block devices looks the same, read blocks yielding processor to other task
- Kernel buffered writes
 - Completion of out-going transfer decoupled from the application, allowing it to continue
- Explicit close

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.38

What's below the surface ??



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.39

Recall: SYSCALL

syscalls.kernelgrok.com

Linux Syscall Reference

Show 10 entries

#	Name	eax	ecx	edx	esi	edi	Definition
0	sys_restart_syscall	0x00	-	-	-	-	kernel/signal.c:2058
1	sys_exit	0x01	int error_code	-	-	-	kernel/exit.c:1046
2	sys_fork	0x02	struct pt_regs *	-	-	-	arch/alpha/kernel/entry.S:716
3	sys_read	0x03	unsigned int fd	char __user *buf	size_t count	-	fs/read_write.c:391
4	sys_write	0x04	unsigned int fd	const char __user *buf	size_t count	-	fs/read_write.c:408
5	sys_open	0x05	const char __user *filename	int flags	int mode	-	fs/open.c:900
6	sys_close	0x06	unsigned int fd	-	-	-	fs/open.c:969
7	sys_waitpid	0x07	pid_t pid	int __user *stat_addr	int options	-	kernel/exit.c:1771
8	sys_creat	0x08	const char __user *pathname	int mode	-	-	fs/open.c:933
9	sys_link	0x09	const char __user *oldname	const char __user *newname	-	-	fs/namel.c:2520

Showing 1 to 10 of 338 entries

Generated from Linux kernel 2.6.35.4 using Exuberant Ctags, Python, and DataTables.
Project on GitHub. Hosted on GitHub Pages.

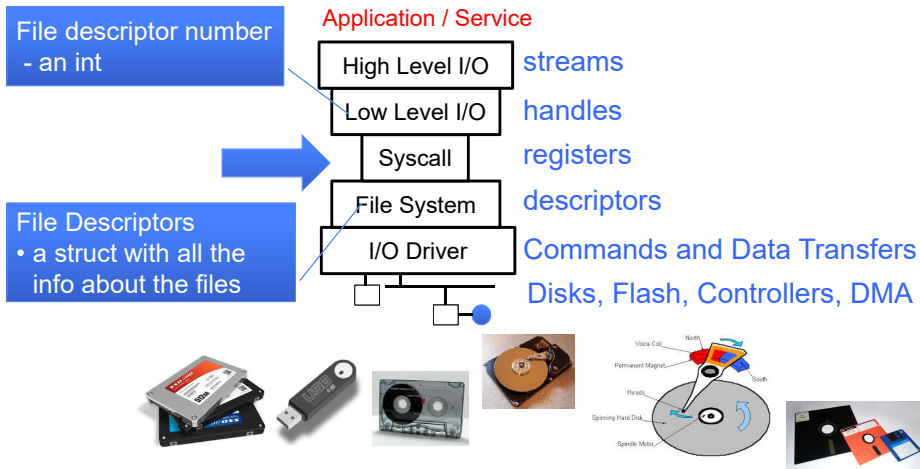
- Low level lib parameters are set up in registers and syscall instruction is issued
 - A type of synchronous exception that enters well-defined entry points into kernel

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.40

What's below the surface ??



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.41

Internal OS File Descriptor

- Internal Data Structure describing everything about the file
 - Where it resides
 - Its status
 - How to access it

- Pointer:
`struct file *file`

```

746 struct file {
747     union {
748         struct list_node fu_list;
749         struct rcu_head fu_rcuhead;
750     } fu;
751     struct path f_path;
752     struct f_dentry f_dentry;
753     struct inode *f_inode; /* caci
754     const struct file_operations *f_op;
755
756     /*
757     * Protects f_ep_links, f_flags.
758     * Must not be taken from IRQ context.
759     */
760     spinlock_t f_lock;
761     atomic_long_t f_count;
762     unsigned int f_flags;
763     fmode_t f_mode;
764     struct mutex f_pos_lock;
765     loff_t f_pos;
766     struct fown_struct f_owner;
767     const struct cred *f_cred;
768     struct file_ra_state f_ra;
769
770     u64 f_version;
771
772 #ifdef CONFIG_SECURITY
773     void *f_security;
774 #endif
775     /* needed for tty driver, and maybe others */
776     void *private_data;
777
778 #ifdef CONFIG_EPOLL
779     /* Used by fs/eventpoll.c to link all the hook:
780     struct list_head f_ep_links;
781     struct list_head f_file_llink;
782 #endif /* #ifdef CONFIG_EPOLL */
783     struct address_space *f_mapping;
784 } __attribute__((aligned(4))); /* test something weird
    
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.42

File System: from syscall to driver

In fs/read_write.c

```

ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EINVAL;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count)))
        return -EINVAL;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
    
```

• Read up to "count" bytes from "file" starting from "pos" into "buf".

• Return error or number of bytes read.

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.43

File System: from syscall to driver

In fs/read_write.c

```

ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count)))
        return -EINVAL;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
    
```

Make sure we are allowed to read this file

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.44

File System: from syscall to driver

In fs/read_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Check if file has read methods

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.45

File System: from syscall to driver

In fs/read_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

•Check whether we can write to buf (e.g., buf is in the user space range)
•unlikely(): hint to branch prediction this condition is unlikely

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.46

File System: from syscall to driver

In fs/read_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Check whether we read from a valid range in the file.

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.47

File System: from syscall to driver

In fs/read_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

If driver provide a read function (f_op->read) use it; otherwise use do_sync_read()

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.48

File System: from syscall to driver

In fs/read_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Notify the parent of this file that the file was read
(see <http://www.fieldses.org/~bfields/kernel/vfs.txt>)

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.49

File System: from syscall to driver

In fs/read_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Update the number of bytes read by "current" task (for scheduling purposes)

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.50

File System: from syscall to driver

In fs/read_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Update the number of read syscalls by "current" task (for scheduling purposes)

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.51

Lower Level Driver

- Associated with particular hardware device
- Registers / Unregisters itself with the kernel
- Handler functions for each of the file operations

```
struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*aio_read) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
    ssize_t (*aio_write) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    int (*ioctl) (struct inode *, struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, struct dentry *, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*flock) (struct file *, int, struct file_lock *);
    [...]
};
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.52

Device Drivers

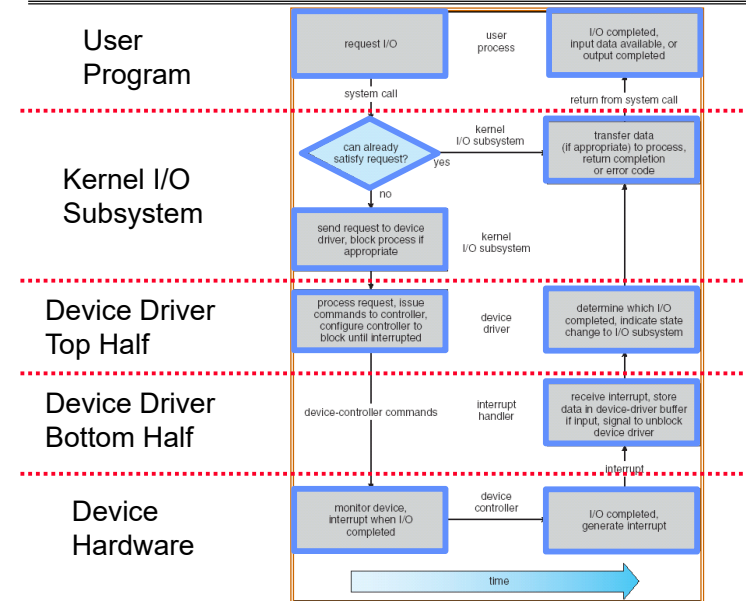
- **Device Driver:** Device-specific code in the kernel that interacts directly with the device hardware
 - Supports a standard, internal interface
 - Same kernel I/O system can interact easily with different device drivers
 - Special device-specific configuration supported with the `ioctl()` system call
- Device Drivers typically divided into two pieces:
 - Top half: accessed in call path from system calls
 - » implements a set of **standard, cross-device calls** like `open()`, `close()`, `read()`, `write()`, `ioctl()`, `strategy()`
 - » This is the kernel's interface to the device driver
 - » Top half will *start* I/O to device, may put thread to sleep until finished
 - Bottom half: run as interrupt routine
 - » Gets input or transfers next block of output
 - » May wake sleeping threads if I/O now complete

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.53

Life Cycle of An I/O Request



2/18/2020


Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.54

Communication between processes

- Can we view files as communication channels?

`write(wfd, wbuf, wlen);`



`n = read(rfd, rbuf, rmax);`

- Producer and Consumer of a file may be distinct processes
 - May be separated in time (or not)
- However, what if data written once and consumed once?
 - Don't we want something more like a queue?
 - Can still look like File I/O!


2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.55

Communication Across the world looks like file IO

`write(wfd, wbuf, wlen);`



`n = read(rfd, rbuf, rmax);`

- Connected queues over the Internet
 - But what's the analog of `open`?
 - What is the namespace?
 - How are they connected in time?

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

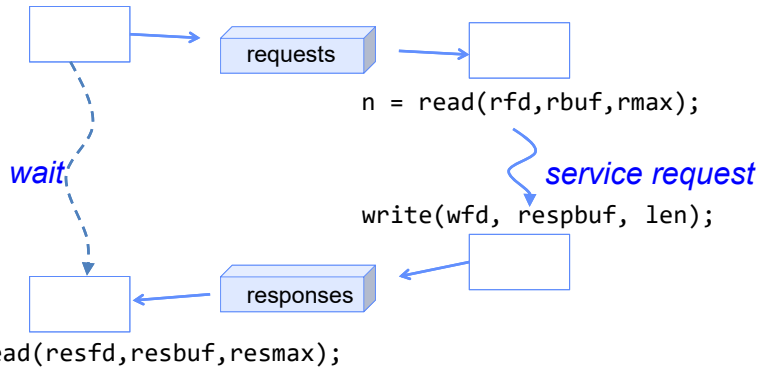
Lec 8.56

Request Response Protocol

Client (issues requests)

Server (performs operations)

```
write(rqfd, rqbuf, buflen);
```



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

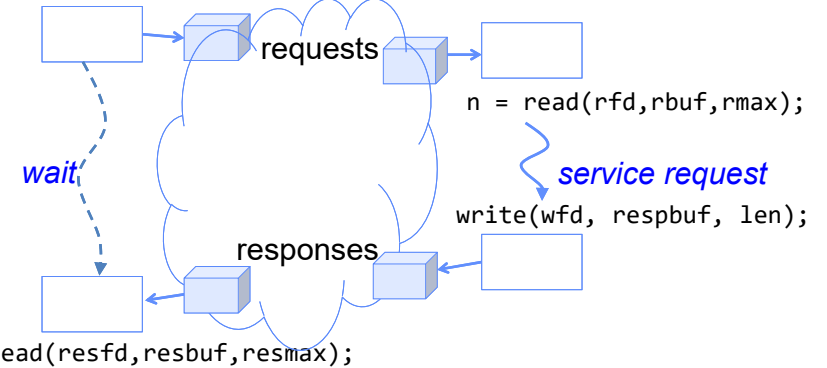
Lec 8.57

Request Response Protocol

Client (issues requests)

Server (performs operations)

```
write(rqfd, rqbuf, buflen);
```




2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.58

Client-Server Models



- File servers, web, FTP, Databases, ...
- Many clients accessing a common server


2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.59

Client-Server Communication

- Client “sometimes on”
 - Initiates a request to the server when interested
 - E.g., Web browser on your laptop or cell phone
 - Doesn’t communicate directly with other clients
 - Needs to know the server’s address
- Server is “always on”
 - Services requests from many client hosts
 - E.g., Web server for the *www.cnn.com* Web site
 - Doesn’t initiate contact with the clients
 - Needs a fixed, well-known address



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.60

Sockets

- **Socket**: an abstraction of a network I/O queue
 - Mechanism for inter-process communication
 - Embodies one side of a communication channel
 - » Same interface regardless of location of other end
 - » Could be local machine (called “UNIX socket”) or remote machine (called “network socket”)
 - First introduced in 4.2 BSD UNIX: big innovation at time
 - » Now most operating systems provide some notion of socket
- Data transfer like files
 - Read / Write against a descriptor
- Over ANY kind of network
 - Local to a machine
 - Over the internet (TCP/IP, UDP/IP)
 - OSI, Appletalk, SNA, IPX, SIP, NS, ...

2/18/2020

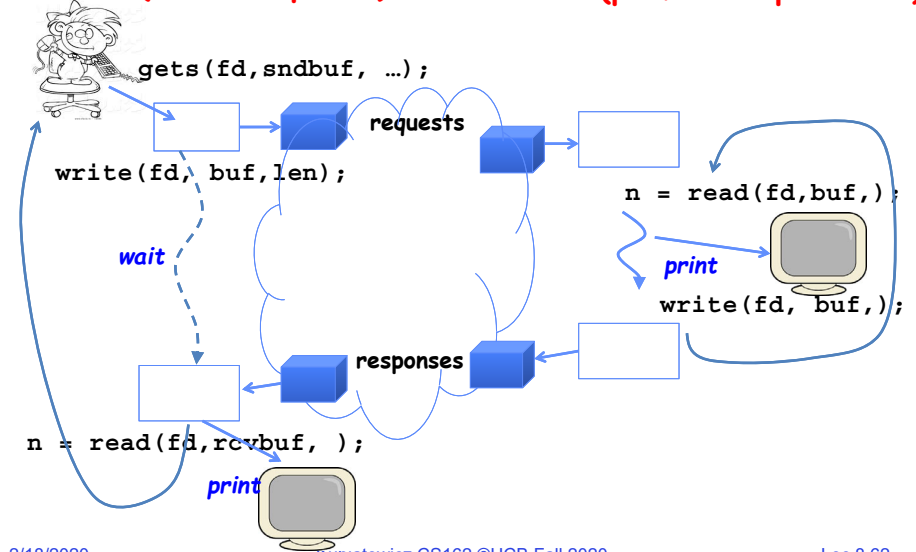
Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.61

Silly Echo Server – running example

Client (issues requests)

Server (performs operations)



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.62

Echo client-server example

```
void client(int sockfd) {
    int n;
    char sndbuf[MAXIN]; char rcvbuf[MAXOUT];
    getreq(sndbuf, MAXIN); /* prompt */
    while (strlen(sndbuf) > 0) {
        write(sockfd, sndbuf, strlen(sndbuf)); /* send */
        memset(rcvbuf, 0, MAXOUT); /* clear */
        n = read(sockfd, rcvbuf, MAXOUT-1); /* receive */
        write(STDOUT_FILENO, rcvbuf, n); /* echo */
        getreq(sndbuf, MAXIN); /* prompt */
    }
}
```

```
void server(int consockfd) {
    char reqbuf[MAXREQ];
    int n;
    while (1) {
        memset(reqbuf, 0, MAXREQ);
        n = read(consockfd, reqbuf, MAXREQ-1); /* Recv */
        if (n <= 0) return;
        n = write(STDOUT_FILENO, reqbuf, strlen(reqbuf));
        n = write(consockfd, reqbuf, strlen(reqbuf)); /* echo */
    }
}
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.63

What assumptions are we making?

- Reliable
 - Write to a file => Read it back. Nothing is lost.
 - Write to a (TCP) socket => Read from the other side, same.
 - Like pipes
- In order (sequential stream)
 - Write X then write Y => read gets X then read gets Y
- When ready?
 - File read gets whatever is there at the time. Assumes writing already took place.
 - Like pipes!

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.64

Socket creation and connection

- File systems provide a collection of permanent objects in structured name space
 - Processes open, read/write/close them
 - Files exist independent of the processes
- Sockets provide a means for processes to communicate (transfer data) to other processes.
- Creation and connection is more complex
- Form 2-way pipes between processes
 - Possibly worlds away
- How do we name them?
- How do these completely independent programs know that the other wants to “talk” to them?

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.65

Namespaces for communication over IP


- Hostname
 - www.eecs.berkeley.edu
- IP address
 - 128.32.244.172 (ipv6?)
- Port Number
 - 0-1023 are “well known” or “system” ports
 - » Superuser privileges to bind to one
 - 1024 – 49151 are “registered” ports (registry)
 - » Assigned by IANA for specific services
 - 49152–65535 ($2^{15}+2^{14}$ to $2^{16}-1$) are “dynamic” or “private”
 - » Automatically allocated as “ephemeral Ports”

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.66

Socket Setup over TCP/IP




- Special kind of socket: **server socket**
 - Has file descriptor
 - Can't read or write
- Two operations:
 1. **listen()**: Start allowing clients to connect
 2. **accept()**: Create a *new socket* for a *particular* client connection

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.67

Socket Setup over TCP/IP



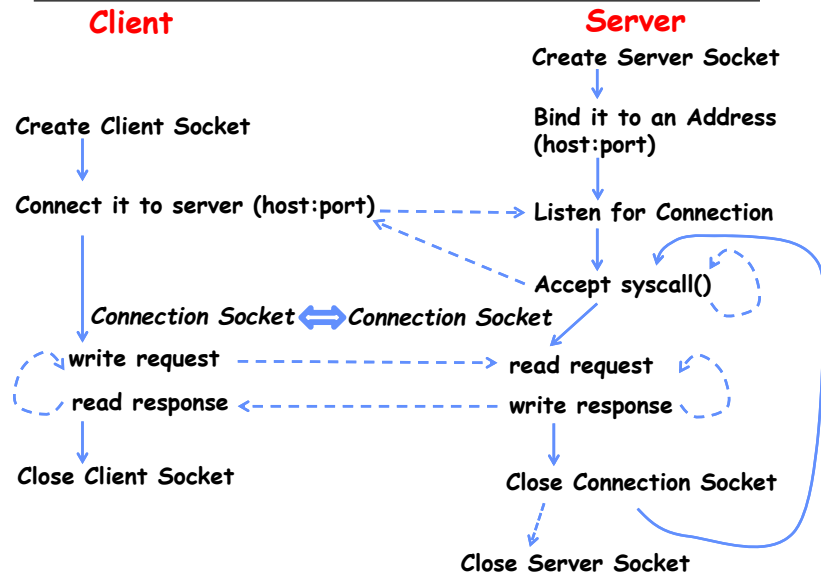
- Server Socket: Listens for new connections
 - Produces new sockets for each unique connection
 - 3-way handshake to establish new connection!
- Things to remember:
 - Connection involves 5 values:
[Client Addr, Client Port, Server Addr, Server Port, Protocol]
 - Often, Client Port “randomly” assigned
 - » Done by OS during client socket setup
 - Server Port often “well known”
 - » 80 (web), 443 (secure web), 25 (sendmail), etc
 - » Well-known ports from 0—1023

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.68

Web Server using Sockets (in concept)



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.69

Client Protocol

```
char *host_name, port_name;

// Create a socket
struct addrinfo *server = lookup_host(host_name, port_name);
int sock_fd = socket(server->ai_family, server->ai_socktype,
                     server->ai_protocol);

// Connect to specified host and port
connect(sock_fd, server->ai_addr, server->ai_addrlen);

// Carry out Client-Server protocol
run_client(sock_fd);

/* Clean up on termination */
close(sock_fd);
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.70

Client: getting the server address

```
struct addrinfo *lookup_host(char *host_name, char *port) {
    struct addrinfo *server;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(hints));
    hints.ai_family = AF_UNSPEC;
    hints.ai_socktype = SOCK_STREAM;

    int rv = getaddrinfo(host_name, port_name,
                        &hints, &server);

    if (rv != 0) {
        printf("getaddrinfo failed: %s\n", gai_strerror(rv));
        return NULL;
    }
    return server;
}
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.71

Server Protocol (v1)

```
// Create socket to listen for client connections
char *port_name;
struct addrinfo *server = setup_address(port_name);
int server_socket = socket(server->ai_family,
                          server->ai_socktype, server->ai_protocol);

// Bind socket to specific port
bind(server_socket, server->ai_addr, server->ai_addrlen);

// Start listening for new client connections
listen(server_socket, MAX_QUEUE);

while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    serve_client(conn_socket);
    close(conn_socket);
}

close(server_socket);
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.72

Server Address - itself

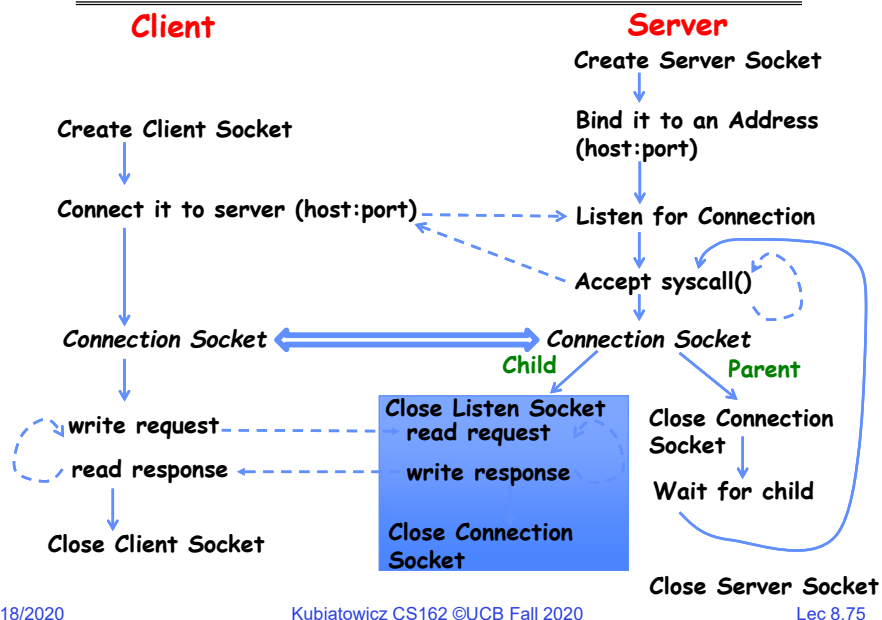
```
struct addrinfo *setup_address(char *port) {
    struct addrinfo *server;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(hints));
    hints.ai_family = AF_UNSPEC;
    hints.ai_socktype = SOCK_STREAM;
    hints.ai_flags = AI_PASSIVE;
    getaddrinfo(NULL, port, &hints, &server);
    return server;
}
```

- Simple form
- Internet Protocol, TCP
- Accepting any connections on the specified port

How does the server protect itself?

- Isolate the handling of each connection
- By forking it off as another process

Sockets With Protection



Server Protocol (v2)

```
// Start listening for new client connections
listen(server_socket, MAX_QUEUE);
while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);

    pid_t pid = fork();
    if (pid == 0) {
        // New process for connection
        // Child process
        close(server_socket);
        serve_client(conn_socket);
        close(conn_socket);
        exit(EXIT_SUCCESS);
    } else {
        // Parent process
        close(conn_socket);
        wait(NULL);
        // Don't need client socket
        // Wait for our (one) child
    }
}
close(server_socket);
```

Concurrent Server

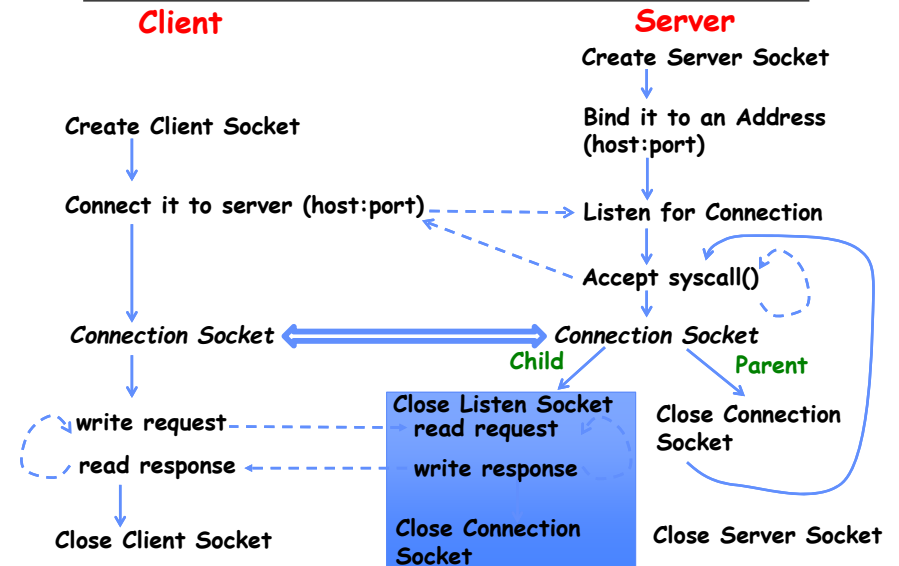
- Listen will queue requests
- Buffering present elsewhere
- But server waits for each connection to terminate before initiating the next

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.77

Sockets With Protection and Parallelism



2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.78

Server Protocol (v3)

```
// Start listening for new client connections
listen(server_socket, MAX_QUEUE);
signal(SIGCHLD, SIG_IGN); // Prevent zombie children
while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);

    pid_t pid = fork(); // New process for connection
    if (pid == 0) {    // Child process
        close(server_socket); // Doesn't need server_socket
        serve_client(conn_socket); // Serve up content to client
        close(conn_socket); // Done with client!
        exit(EXIT_SUCCESS);
    } else {           // Parent process
        close(conn_socket); // Don't need client socket
        // wait(NULL); // Don't wait (SIGCHLD ignored, above)
    }
}
close(server_socket);
```

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.79

Conclusion (I)

- System Call Interface is “narrow waist” between user programs and kernel
- Streaming IO: modeled as a stream of bytes
 - Most streaming I/O functions start with “f” (like “fread”)
 - Data buffered automatically by C-library functions
- Low-level I/O:
 - File descriptors are integers
 - Low-level I/O supported directly at system call level
- STDIN / STDOUT enable composition in Unix
 - Use of pipe symbols connects STDOUT and STDIN
 - » find | grep | wc ...

2/18/2020

Kubiatowicz CS162 ©UCB Fall 2020

Lec 8.80

Conclusion (II)

- Device Driver: Device-specific code in the kernel that interacts directly with the device hardware
 - Supports a standard, internal interface
 - Same kernel I/O system can interact easily with different device drivers
- File abstraction works for inter-processes communication (local or Internet)
- Socket: an abstraction of a network I/O queue
 - Mechanism for inter-process communication